

hw1

February 3, 2015

1 Homework 1

- Name: Austin
- SID: 23826762
- Repo: Open up hw1.ipynb in IPython Notebook.

```
In [1]: %matplotlib inline
import matplotlib.pyplot as plt
import scipy.io
import numpy
from sklearn import svm

# from http://stackoverflow.com/q/4601373/1222351
def permute_both(a, b):
    p = numpy.random.permutation(len(a))
    return a[p], b[p]

In [2]: # Load the data
mat = scipy.io.loadmat('data/digit-dataset/train.mat')
images = numpy.reshape(mat["train_images"], (1, -1, 60000))[0].T
labels = mat["train_labels"]

In [3]: # Shuffle the data in parallel
images, labels = permute_both(images, labels)
labels = numpy.ravel(labels)
```

1.1 Problems 1 & 2

The accuracies approach 80% as the number of training examples increase.

From the confusion matrix, it appears that 5s and 8s are the most commonly misidentified digits, often labeled as the reverse. Meanwhile, 0s and 1s are very accurately identified.

```
In [4]: i = 0
x, y = [100, 200, 500, 1000, 2000, 5000, 10000], []
for size in x:
    s = svm.LinearSVC()
    s.fit(images[i:i + size], labels[i: i + size])
    score = s.score(images[-10000:], labels[-10000:])
    y.append(score)
    print(size, " - ", score)
    i += size + 1


# Make a confusion matrix
from sklearn.metrics import confusion_matrix
```

```


cm = confusion_matrix(s.predict(images[-10000:]), labels[-10000:])
plt.matshow(cm)
plt.title('Confusion matrix')
plt.colorbar()
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()

```


100 - 0.6429




200 - 0.6908




500 - 0.7459




1000 - 0.7503




2000 - 0.7418



5000 - 0.7776




10000 - 0.8025



In [5]: plt.plot(x, y)

Out[5]: [<matplotlib.lines.Line2D at 0xadf4d30>]



1.2 Problem 3

Cross-validation allows you to train and test against the same set of data, with an overall effect of improving accuracy by allowing more training data to be used.

The optimal 'C' I've found is $5 * 10^{-7}$, providing a cross-validation accuracy of 0.86. To converge on this value, I first started C on 10^{-10} , then repeatedly multiply by a factor of 100 up to 10^{10} , to get an order of magnitude approximation. Then I tried increasing C with a linear step size.

```
In [ ]: # Reshuffle the data
        images, labels = permute_both(images, labels)
        labels = numpy.ravel(labels)

        c = 1e-7

        for _ in range(5):

            # 10-fold cross validation
            total = 0
            for k in range(10):
                kimages, klabels = [], []
                s = svm.LinearSVC(C=c)
                for i in range(10):
                    if i != k:
                        start, end = i * 1000, i * 1000 + 1000
                        kimages.append(images[start:end])
                        klabels.append(labels[start:end])
                s.fit(numpy.concatenate(kimages), numpy.concatenate(klabels))

                start, end = k * 1000, k * 1000 + 1000
                total += s.score(images[start:end], labels[start:end])
            print(c, "-", total / (k + 1))

        c += 2e-7
```

1.3 Kaggle - Digits

With $C = 5 * 10^{-7}$ on a LinearSVC, I got a Kaggle score of 0.87820.

```
In [4]: # Load the data
        mat = scipy.io.loadmat('data/digit-dataset/train.mat')
        images = numpy.reshape(mat["train_images"], (1, -1, 60000))[0].T
        labels = mat["train_labels"]
        labels = numpy.ravel(labels)

        # Train the SVM on all the training data
        s = svm.LinearSVC(C=1e-7).fit(images, labels)

        # Predict the labels for the test data
        mat = scipy.io.loadmat('data/digit-dataset/test.mat')
        test = numpy.reshape(mat["test_images"], (1, -1, 10000))[0].T
        result = s.predict(test)

In [16]: # Write the results to a csv
         f = open('digits.csv', 'w')
```

```
f.write('Id,Category\n')
for i in range(len(result)):
    f.write("{0},{1}\n".format(i + 1, result[i]))
f.close()
```

1.4 Problem 4

The optimal 'C' I've found is 90, providing a cross-validation accuracy of 0.813. This was derived with the same method – starting C on 10^{-10} , then repeatedly multiply by a factor of 100 up to 10^{10} , to get an order of magnitude approximation, and finally increasing C with a linear step size.

```
In [2]: # Load the data
mat = scipy.io.loadmat('data/spam-dataset/spam_data.mat')
emails = mat['training_data']
labels = numpy.ravel(mat['training_labels'])
emails, labels = permute_both(emails, labels)

In [ ]: # k-fold cross validation to optimize C
emails, labels = permute_both(emails, labels)

folds = 12
interval = 431

c = 90

for _ in range(10):
    # cross validation
    total = 0
    for k in range(folds):
        kemails, klabels = [], []
        s = svm.SVC(C=c)
        for i in range(folds):
            if i != k:
                start, end = i * interval, (i + 1) * interval
                kemails.append(emails[start:end])
                klabels.append(labels[start:end])
        s.fit(numpy.concatenate(kemails), numpy.concatenate(klabels))

        start, end = k * interval, (k + 1) * interval
        total += s.score(emails[start:end], labels[start:end])
    print(c, "-", total / (folds))
    c *= 3
```

1.5 Kaggle - Spam

With $C = 90$ on a LinearSVC, I got a Kaggle score of 0.82548.

```
In [6]: # Load the data
mat = scipy.io.loadmat('data/spam-dataset/spam_data.mat')
emails = mat['training_data']
labels = numpy.ravel(mat['training_labels'])

# Train the SVM on all the training data
s = svm.SVC(C=90).fit(emails, labels)
```

```
# Predict the labels for the test data
result = s.predict(mat['test_data'])

In [7]: # Write the results to a csv
f = open('spam.csv', 'w')
f.write('Id,Category\n')
for i in range(len(result)):
    f.write("{0},{1}\n".format(i + 1, result[i]))
f.close()
```